


Event-level parallelism: attention is AIE-tile-bound (~30×), the full model is PL/AIE-bound (~5×)

Measured throughput (HW, golden-fed)
attention block >> full model



Event-parallel ceiling = min over resources
full model: PL/AIE bind @5–6 • attn block: AIE-tile bind @30



Scale-out throughput
(latency unchanged; data-parallel)

