




Full-model impact of offloading ALL attention to AIE (all-PL vs all-AIE, measured on VCK190)




1.15× faster



+15%



~flat (static-dominated)



full model: MSE 0.521→0.610 (+17%)