

Unquantized Reference: same AIE kernels, two arithmetics, vs PyTorch
(worst of the 6 attention blocks per event, x86sim, retrained weights, 2000 events)

